Fast Principal Component Analysis of Large-Scale Genome-Wide Data

Gad Abraham[1,*], Michael Inouye[1]

1 Medical Systems Biology, Department of Pathology and Department of Microbiology & Immunology, University of Melbourne, Parkville 3010, Victoria, Australia
∗ E-mail: gad.abraham@unimelb.edu.au

Abstract

Principal component analysis (PCA) is routinely used to analyze genome-wide single-nucleotide polymorphism (SNP) data, for detecting population structure and potential outliers. However, the size of SNP datasets has increased immensely in recent years and PCA of large datasets has become a time consuming task. We have developed `flashpca`, a highly efficient PCA implementation based on randomized algorithms, which delivers identical accuracy compared with existing tools in substantially less time. We demonstrate the utility of `flashpca` on both HapMap3 and on a large Immunochip dataset. For the latter, `flashpca` performed PCA of 15,000 individuals up to 125 times faster than existing tools, with identical results, and PCA of 150,000 individuals using `flashpca` completed in 4 hours. The increasing size of SNP datasets will make tools such as `flashpca` essential as traditional approaches will not adequately scale. This approach will also help to scale other applications that leverage PCA or eigen-decomposition to substantially larger datasets.

Introduction

Principal component analysis (PCA) is a widely-used tool in genomics and statistical genetics, employed to infer cryptic population structure from genome-wide data such as single nucleotide polymorphisms (SNPs) [1,2], and/or to identify outlier individuals which may need to be removed prior to further analysis, such as genome-wide association studies (GWAS). This is based on the fact that such population structure can confound SNP-phenotype associations, resulting in some SNPs spuriously being called as associated with the phenotype (false positives). The principal components (PCs) of SNP data have been shown to map well to geographic distances between human populations [1,3], thus capturing the coarse-grain allelic variation between them these groups.

However, traditional approaches to computing the PCA, such as those employed by the popular EIGENSOFT suite [1], are computationally expensive. For example, PCA based on the singular value decomposition (SVD) scales as $\mathcal{O}(\min{(N^2p, Np^2)})$ and for eigen-decomposition it is $\mathcal{O}(\min{(N^3, p^3)})$, where N and p are the number of samples and SNPs, respectively. This makes it time-consuming to perform PCA on large cohorts such as those routinely being analyzed today, involving millions of assayed or imputed SNPs and tens of thousands of individuals, with this difficulty only increasing in the future with the availability of even larger studies.

In recent years, research into randomized matrix algorithms has yielded alternative approaches for performing PCA that are far more computationally tractable while maintaining high accuracy relative to the traditional "exact" algorithms [4]. These algorithms are especially useful when we are interested in finding only the first few principal components (PCs) of the data, as is often the case in genomic analysis.

Here we present `flashpca`, an efficient tool for performing PCA analysis on large genome-wide data, based on randomized algorithms. Our approach is highly efficient, allowing the user to perform PCA on large datasets (100,000 individuals or more), while achieving identical results to traditional methods. In addition, our tool allows the user to "whiten" the data, producing data corrected for the principal components which can then be analyzed using standard approaches such as linear models [5].

Results

First we used an LD-pruned HapMap3 genotype data [6] consisting of 957 human individuals across 11 populations assayed for 14,389 SNPs (Materials). We compared `flashpca` with `smartpca` from EIGEN-SOFT v4.2 [1] and `shellfish`[2]. In addition, we included the R 3.0.2 `prcomp` function which is based on SVD rather than eigen-decomposition, after replacing its original standardization with the one used by `smartpca` (Equation 4). The analysis of HapMap3 data revealed the expected ancestry via the first two PCS (Figure 1a), with individuals of east Asian origin (CHB, CHD, JPT) clustered in the bottom right-hand side corner, those of European origin (TSI, CEU) in the top, and those of African origin in the bottom left-hand side corner (ASW, LWK, YRI, and MKK). All methods showed close to perfect agreement on the top PC (Figure 1b), with an absolute correlation of 1.00 between the 1st PC of each pair of methods (the sign of the eigenvectors is arbitrary hence the correlation may be negative as well). The next nine PCs showed close to perfect agreement as well (see the the online documentation).

Next, we analyzed an LD-pruned celiac disease Immunochip dataset consisting of 16,002 individuals and 43,049 SNPs [7] (Materials). We created an artificial dataset of 160,020 samples by duplicating the original dataset 10 times. We then randomly sampled subsets of the large dataset with increasing size ($N = 500, 1000, 2500, 5000, 7500, 10,000, 15,000$), and recorded wall time for `flashpca`, `smartpca`, and `shellfish` performing PCA on these subsets. Due to the substantial time required by `shellfish` and `smartpca` to complete the largest runs, we also ran larger subsets ($N = 50000, 100,000, 150,000$) using `flashpca` only (we attempted to run `shellfish` on the $N = 50,000$ dataset but it did not complete due to running out of memory, and stopped `smartpca` after 100,000sec). Each experiment was repeated three times. All programs used multi-threaded mode with 8 threads. All experiments were run on a machine with 4×10-core Intel Xeon E7-4850 CPU @ 2.00GHz with 512GiB RAM running 64-bit Ubuntu Linux 12.04.

Figure 2 shows that `flashpca` was substantially faster than either `smartpca` or `shellfish`: for analysis of 15,000 samples, `flashpca` took an average of 8 minutes, whereas `smartpca` required an average of almost 17h ($\times 125$ slower). Examining the large subsets, `flashpca` was able to analyze a dataset of $N = 150,000$ in ~ 4h, which would not be sufficient time for `smartpca` to complete a PCA on 10,000 samples, as 6.5h were required for that. While `shellfish` was also substantially faster than `smartpca`, it was still considerably slower than `flashpca` when the number of individuals was large, with a run time of $\sim$ 1h for $N = 15,000$ (and did not complete for subsets with $N \geq 50,000$).

Importantly, the time taken by `flashpca` to compute PCA on $N = 150,000$ does not differ much from the time taken on the $N = 100,000$ subset; this is because `flashpca` automatically transposes the data when $N > p$, as the PCA can be computed on the original data or its transpose with only minor modifications to the algorithm; computing the $p \times p$ covariance matrix and its eigen-decomposition has lower computational complexity than using the $N \times N$ covariance matrix, for the same values of N and p, hence when $N > p$ the main computational cost will not grow substantially with N.

Discussion

Principal component analysis is an important tool in genomics for discovery of population structure or other latent structure in the data, such as batch effects. Early approaches such as `smartpca` from EIGENSOFT have proven useful for this goal and have been widely used for analysis of SNP datasets. However, many current datasets assay tens of thousands of individuals, making traditional approaches extremely time consuming. In contrast, our approach, `flashpca`, is based on careful combination of randomized algorithms for PCA together with parallelization, and allows the analyst to easily perform

---

[1]`http://www.hsph.harvard.edu/alkes-price/software/`
[2]`http://www.stats.ox.ac.uk/~davison/software/shellfish/shellfish.php`

PCA on large datasets consisting of many thousands of individuals in a matter of minutes to hours and with no loss in accuracy.

More generally, the approach behind `flashpca` could prove useful for accelerating other methods that depend on performing a large number of eigen-decompositions across many samples, such as varLD [8] which assesses local differences in SNP LD between populations or FaST-LMM which implements linear mixed models of SNP data [9].

Materials and Methods

Datasets

HapMap3

The HapMap3 phase 3 dataset consists of 1184 human individuals across 11 populations (ASW: African ancestry in Southwest USA; CEU: Utah residents with Northern and Western European ancestry from the CEPH collection; CHB: Han Chinese in Beijing, China; CHD: Chinese in Metropolitan Denver, Colorado; GIH: Gujarati Indians in Houston, Texas; JPT: Japanese in Tokyo, Japan; LWK: Luhya in Webuye, Kenya; MEX: Mexican ancestry in Los Angeles, California; MKK: Maasai in Kinyawa, Kenya; TSI: Toscani in Italia; YRI: Yoruba in Ibadan, Nigeria) assayed for 1,440,616 SNPs [6]. We performed QC on the data, including removal of SNPs with MAF$< 1\%$, missingness $> 1\%$, and deviation from Hardy-Weinberg equilibrium $P < 5 \times 10^{-6}$. We removed non-founders and individuals with genotyping missingness $> 1\%$, leaving 957 individuals. Next, we removed several regions of high LD and/or known inversions (chr5: 44Mb–51.5Mb, chr6: 25Mb–33.5Mb, chr8: 8Mb–12Mb, chr11: 45Mb–57Mb) [10]. Finally, we used PLINK [11] `--indep-pairwise 1000 10 0.02` to thin the SNPs by LD ($r^2 < 0.02$), leaving 14,389 SNPs.

Celiac Disease Immunochip

The celiac disease Immunochip dataset [7] consists of 16,002 case/control individuals of British descent, assayed for 139,553 SNPs. The QC has been previously described [7]. In addition, we removed SNPs with MAF $< 0.5\%$ and non-autosomal SNPs, leaving 115,746 SNPs. Next, we removed the same four regions as for HapMap3, and finally, we thinned the SNPs by LD with PLINK `--indep-pairwise 50 10 0.5`, leaving 43,049 SNPs.

Principal Component Analysis

We represent the genotypes as an $N \times p$ matrix $\mathbf{X}$, where the N samples index the rows and the p SNPs index the columns.

PCA relies on finding the eigenvectors of the $N \times N$ covariance matrix $\mathbf{\Sigma} = \mathbf{X}\mathbf{X}^T$ (for notational clarity we will ignore the scaling $\frac{1}{N-1}$ factor which can be accounted for later). This decomposition is performed using either the singular value decomposition (SVD) of the matrix $\mathbf{X}$ or the eigen-decomposition of the covariance matrix itself.

The SVD of $\mathbf{X}$ is

$$\mathbf{X} = \mathbf{U}\mathbf{D}\mathbf{V}^T, \tag{1}$$

where $\mathbf{U}$ an is an $N \times k$ orthogonal matrix ($\mathbf{U}^T\mathbf{U} = \mathbf{I}$) the columns of which are the eigenvectors of $\mathbf{X}\mathbf{X}^T$, $\mathbf{D}$ is a $k \times k$ diagonal matrix of singular values (square-root of the eigenvalues of $\mathbf{X}\mathbf{X}^T$ and $\mathbf{X}^T\mathbf{X}$), and $\mathbf{V}$ is a $p \times k$ orthogonal matrix ($\mathbf{V}^T\mathbf{V} = \mathbf{I}$) of the eigenvectors of $\mathbf{X}^T\mathbf{X}$, where k is the matrix rank (this SVD is also called the "economy SVD"). Note that SVD does not require the covariance matrix $\mathbf{\Sigma}$ to be computed explicitly.

In the eigen-decomposition approach, the covariance matrix $\mathbf{\Sigma}$ is first explicitly computed, then the eigen-decomposition is performed such that

$$\mathbf{\Sigma} = \mathbf{U\Lambda U}^T, \tag{2}$$

where $\text{diag}(\mathbf{\Lambda}) = \lambda_1, \ldots, \lambda_k$ are the eigenvalues and $\mathbf{U}$ is the matrix of eigenvectors as before.

The principal components (PCs) $\mathbf{P}$ of the data are given by the projection of the data onto the eigenvectors

$$\mathbf{P} = \mathbf{XV} = \mathbf{UD}, \tag{3}$$

where we usually truncate the matrix $\mathbf{P}$ to have as many columns as required for any down-stream analysis (say, 10).

Note that some tools, such as `smartpca` and `shellfish`, output the eigenvectors $\mathbf{U}$ as the principal components without weighting by the singular values $\mathbf{D}$, leading to different scales for the PCs. In addition, since the (empirical) covariance is typically scaled by a factor of $\frac{1}{N-1}$, then in order to maintain the interpretation of the singular values $\mathbf{D}$ as the square-root of the eigen-values of the scaled covariance $\frac{1}{N-1}\mathbf{XX}^T$, the singular values must be scaled by a factor of $\frac{1}{\sqrt{N-1}}$ as well (as implemented in R's `prcomp`). Note, however, that these scale differences have no effect on the interpretation of the principal components for ascertaining or correcting for potential population structure in data.

In traditional PCA, such as that implemented in R's `prcomp`, prior to running the SVD/eigen-decomposition itself, the matrix $\mathbf{X}$ is first mean-centered by subtracting the per-column (SNP) average from each column. In contrast, EIGENSOFT [1], first centers the data, then divides by a quantity proportional to the standard deviation of a binomial random variable with success probability p_j

$$\tilde{\mathbf{X}}_{ij} = \frac{\mathbf{X}_{ij} - \mu_j}{\sqrt{p_j(1 - p_j)}}, \tag{4}$$

where $p_j = \mu_j/2$ and $\mu_j = \frac{1}{N}\sum_{i=1}^{N} \mathbf{X}_{ij}$. To maintain compatibility with `smartpca`, `flashpca` employs the same standardization method by default.

Fast Principal Component Analysis

Performing PCA on large matrices (with both N and p large), is time consuming with traditional approaches. To enable fast PCA, we employ an algorithm based on a randomized PCA approach [4]. Briefly, randomized PCA relies on first constructing a relatively small matrix that captures the top eigenvalues and eigenvectors of the original data, with high probability. Next, standard SVD or eigen-decomposition is performed on this reduced matrix, producing near identical results to what would have been achieved using a full analysis of the original data. Since in most genomic applications we are interested only in a few of the top eigenvectors of the data (typically 10), this allows reduction of the data to a substantially smaller matrix, and the computational cost of decomposing this matrix is negligible (see Algorithm 1). (Note, however, that this method is general and is as useful for extracting a much larger number of PCs).

Focusing on the eigen-decomposition approach (Equation 2), the two main computational bottlenecks are (i) computing the $N \times N$ covariance matrix $\mathbf{\Sigma}$ and (ii) when N is large, the eigen-decomposition step itself. In our fast PCA approach, the first bottleneck cannot be avoided but can be mitigated through parallel computation (see Implementation). The second bottleneck is circumvented via the randomized approach, by constructing a matrix $\mathbf{B}$ of size $(d + k) \times p$, from which the small $(d + k) \times (d + k)$ matrix $\mathbf{BB}^T$ is formed and used for the eigen-decomposition, where d is the number of required eigenvectors (say 10), and k is the number of auxiliary dimensions used to increase accuracy which can be discarded later. We have found that a total dimension of 200 is more than sufficient for producing good results in the first 10 PCs; hence the eigen-decomposition need only be performed on 200×200 matrix while producing near identical results to a full PCA on the original data.

Algorithm 1: Pseudocode for the eigen-decomposition variant of the fast PCA, based on the randomized algorithm of [4] for the case where $N < p$. standardize$(\cdot)$ is the standardization in Equation 4. randn(m, n) is a function generating an $m \times n$ iid multivariate normal matrix, k is the user-defined number of extra dimensions, qr$(\cdot)$ is the QR decomposition, normalize$(\cdot)$ is a function that divides each column j by its ℓ_2 norm $\sqrt{\sum_{i=1}^{N} \mathbf{X}_{ij}^2}$, eigen$(\cdot)$ is the eigen-decomposition producing the eigenvectors $\tilde{\mathbf{U}}$ and vector of eigenvalues $\boldsymbol{\lambda}$.

$\mathbf{M} = \text{standardize}(\mathbf{X})$
$\mathbf{R} = \text{randn}(p, d + k)$
$\mathbf{Y} = \text{normalize}(\mathbf{MR})$
$\boldsymbol{\Sigma} = \mathbf{MM}^T$
for *iter = 1:maxiter* **do**
$\quad | \quad \mathbf{Y} = \text{normalize}(\boldsymbol{\Sigma}\mathbf{Y}))$
end
$[\mathbf{Q}, \mathbf{R}] = \text{qr}(\mathbf{Y})$
$\mathbf{B} = \mathbf{Q}^T\mathbf{M}$
$\mathbf{S} = \mathbf{BB}^T$
$[\tilde{\mathbf{U}}, \boldsymbol{\lambda}] = \text{eigen}(\mathbf{S})$
$\mathbf{U} = \mathbf{Q}\tilde{\mathbf{U}}_{\cdot,1:d}$
$\mathbf{D} = \text{diag}(\sqrt{\frac{\boldsymbol{\lambda}}{N-1}})$
$\mathbf{P} = \mathbf{UD}$

Another computational shortcut for the case where $N > p$ is simply transposing the data $\mathbf{X}$, then standardizing the rows instead of the columns. An identical PCA algorithm is then run on the transposed data, with the only difference being that the estimated matrix $\mathbf{U}$ will now contain the eigenvectors of $\mathbf{X}^T\mathbf{X}$ instead of $\mathbf{XX}^T$, and hence the final principal components will be $\mathbf{P} = \mathbf{XU}$. This procedure makes it possible to analyze large datasets with $N \gg p$ at a cost not much greater than when $N = p$.

While the SVD approach is generally recommended for reasons of better numerical stability and speed, we have found that when N and p are in the thousands, the above eigen-decomposition approach is substantially faster since the decomposition is performed on the small matrix $\mathbf{S} = \mathbf{BB}^T$ which is of size $(d+k) \times (d+k)$ rather than a more expensive SVD of the matrix $\mathbf{B}$, which is an $N \times (d+k)$ matrix, with no discernible effect on accuracy of the top principal components; however, both methods are available in `flashpca`.

PCA and ZCA Whitening

Whitening is rotation of data such that the covariance is the identity matrix $\mathbf{I}$ (or some scalar multiple of it). In our case, we are interested in whitening the samples to remove the effects of population structure or relatedness, making the samples unrelated and the sample covariance matrix $\mathbf{XX}^T = \mathbf{I}$. The PCA whitened matrix $\mathbf{X}_{\text{PCA}}$ is given by

$$\mathbf{X}_{\text{PCA}} = \mathbf{D}^{-1}\mathbf{U}^T\mathbf{X}, \tag{5}$$

and the so-called ZCA (zero phase component analysis) whitened matrix is

$$\mathbf{X}_{\text{ZCA}} = \mathbf{U}\mathbf{X}_{\text{PCA}}. \tag{6}$$

The difference between the PCA and ZCA whitening matrices is that the former may have a smaller number of rows than $\mathbf{X}$ (and equal to the total number of eigenvalues extracted from the randomized PCA) whereas the ZCA whitening matrix has the full number of rows as in the original matrix $\mathbf{X}$, and hence the ZCA matrix is more useful for further downstream analysis replacing the original data $\mathbf{X}$.

After whitening the genotypes, we may whiten any quantitative phenotypes, represented here as an $N \times K$ matrix $\mathbf{Y}$ using the eigenvalues/eigenvectors

$$\mathbf{Y}_{\text{ZCA}} = \mathbf{U}\mathbf{D}^{-1}\mathbf{U}^T\mathbf{Y}. \tag{7}$$

This procedure is useful for quantitative trait association studies where we wish to remove the population structure component prior to performing an association analysis [5], and is conceptually related to generalized least squares (GLS).

Note that the whitened data are no longer representable as integer allele dosages but are real numbers and must be analyzed as such, and that the space requirements for storing whitened data are substantially higher than for genotypes.

Implementation

`flashpca` is implemented in C++ and relies on Eigen [12], a C++ header-only library of numerical and linear algebra algorithms, which allows for native parallelization of certain computations through OpenMP threads when multiple CPU cores are available. `flashpca` natively reads PLINK [11] SNP-major BED files, avoiding the need to convert these files to other formats. `flashpca` is licensed under the GNU Public License (GPL) v3; source code and documentation are available at `https://github.com/gabraham/flashpca`.

Prior to PCA, thinning of the SNPs by LD and removal of regions with known high LD or other artefacts such as inversions have been recommended [1, 10], as high correlations between the SNPs can distort the resulting eigenvectors. For this purpose we recommend using PLINK v2 [3] which is substantially faster than PLINK v1.07. Reducing a dataset to $\sim$ 10,000–50,000 SNPs is usually sufficient to achieve an accurate PCA, and can be done using `--indep-pairwise`.

Acknowledgments

References

1. Price AL, Patterson NJ, Plenge RM, Weinblatt ME, et al. (2006) Principal components analysis corrects for stratification in genome-wide association studies. Nat Genet 38: 904–909.

2. Patterson N, Price AL, Reich D (2006) Population Structure and Eigenanalysis. PLoS Genet 2: e190.

3. Novembre J, Johnson T, Bryc K, Kutalik Z, Boyko AR, et al. (2008) Genes mirror geography within Europe. Nature 456: 98–101.

4. Halko N, Martinsson PG, Shkolnisky Y, Tygert M (2011) An Algorithm for the Principal Component Analysis of Large Data Sets. SIAM Journal on Scientific Computing 33: 2580–2594.

5. Rakitsch B, Lippert C, Stegle O, Borgwardt K (2013) A Lasso Multi-Marker Mixed Model for Association Mapping with Population Structure Correction. Bioinformatics 2: 206–214.

---

[3] `https://www.cog-genomics.org/plink2`

6. International HapMap 3 Consortium (2010) Integrating common and rare genetic variation in diverse human populations. Nature 467: 52–58.

7. Trynka G, Hunt Ka, Bockett Na, Romanos J, Mistry V, et al. (2011) Dense genotyping identifies and localizes multiple common and rare variant association signals in celiac disease. Nat Genet 43: 1193–201.

8. Ong RTH, Teo YY (2010) varLD: a program for quantifying variation in linkage disequilibrium patterns between populations. Bioinformatics 26: 1269–70.

9. Lippert C, Listgarten J, Liu Y, Kadie CM, Davidson RI, et al. (2011) FaST linear mixed models for genome-wide association studies. Nature Methods 8: 833–5.

10. Fellay J, Ge D, Shianna K, Colombo S, Ledergerber B, et al. (2009) Common Genetic Variation and the Control of HIV-1 in Humans. PLoS Genet 5: e1000791.

11. Purcell S, Neale B, Todd-Brown K, Thomas L, Ferreira MA, et al. (2007) PLINK: a tool set for whole-genome association and population-based linkage analyses. Am J Hum Genet 81: 559–575.

12. Guennebaud G, Jacob B, et al. (2010). Eigen v3. http://eigen.tuxfamily.org.

Figure Legends

Figure 1: (a) The first two principal components from analyzing the HapMap3 dataset. (b) Scatter plots showing near-perfect absolute Pearson correlation (lower left-hand corner) between the 1st PCs estimated by `smartpa`, `flashpca`, `shellfish`, and R's `prcomp` (using the standardization from Equation 4). Note that since eigenvectors are only defined up to sign, the correlations may be negative as well as positive. In addition, the scale of the PCs may differ between the methods, however this has no bearing on the interpretation of the PCs.

Figure 2: Wall time (seconds) for `flashpca` versus EIGENSOFT's `smartpca` and `shellfish` on increasing subsets of the celiac disease dataset, employing multi-threaded mode (8 threads), using 43,049 SNPs. `shellfish` did not complete PCA for the $N \geq 50,000$ subsets, and `smartpca` was stopped after 100,000sec. The results shown are averages over 3 runs.

Figure 1

Figure 2